

# Fast and Accurate Image segmentation for medical imaging data

Team 17  
Xincheng Tan, Yixin Lei, Hanwen Zhang, Scarlett Gong

00

# Objective

Parallel Fast and Accurate  
Image Segmentation on Cell Nucleus

# Image Segmentation of Cell Nuclei

Data source: 2018 Data Science Bowl challenge

**670+**

## Training set

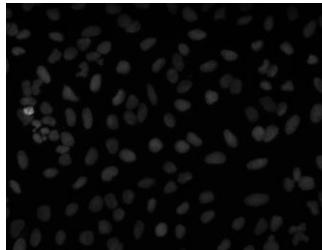
Images of different cell types and tissues

**~65**

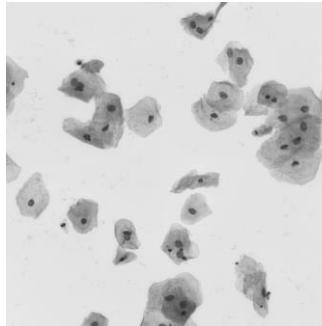
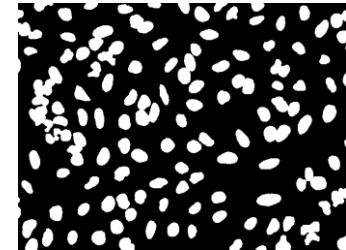
## Test set

Without ground truth masks

Original image



Ground Truth Masks



# Application Type

An **offline, compute intensive toolbox for nuclei segmentation which**

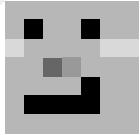
- Provides **three** segmentation models:
  - **Otsu's Method** (threshold method implemented in C++)
  - **Boruvka's Algorithm** (tree-based method implemented in C++)
  - **UNet** (Keras + Tensorflow deep learning in Python)
- Allows fast nuclei segmentation for new images
- **Big Compute + Big Data** (more on the former)
- **Cloud infrastructure:** Google Cloud (VM and cluster)

01

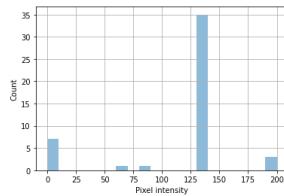
# Otsu's method

# Otsu's Method Algorithm overview

Read image



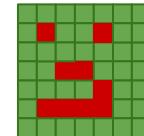
Pixel intensity hist.



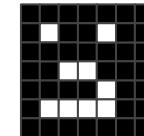
Find threshold

0, 0, 0, 0, 0, 0, 0, 0, ...  
60, .....  
87, .....  
130, 130, 130, 130, 130, 130, ...  
200, 200, 200

Final segmentation



Write result



**Step 0:**

Read in image

**Step 1:**

Build the histogram  
of pixel intensity.

$O(N)$

**Step 2:**

Exhaustive search of  
threshold to minimizes  
inter-class variance.  
 $O(256^2 * 256) \sim O(1)$

**Step 3:**

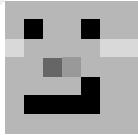
Final image  
segmentation based  
on the threshold.  
 $O(N)$  complexity

**Step 4:**

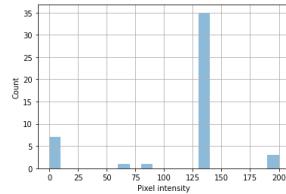
Write out image  
result

# Otsu's Method Algorithm overview

Read image



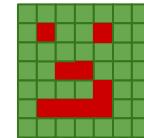
Pixel intensity hist.



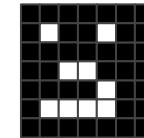
Find threshold

0, 0, 0, 0, 0, 0, 0, 0, ...  
60, .....  
87, .....  
130, 130, 130, 130, 130, 130, ...  
200, 200, 200

Final segmentation



Write result



**Step 0:**

Read in image

**Step 1:**

Build the histogram of pixel intensity.

$O(N)$

**Step 2:**

Exhaustive search of threshold to minimizes inter-class variance.

$O(256 \times 256) \sim O(1)$

**Step 3:**

Final image segmentation based on the threshold.

$O(N)$  complexity

**Step 4:**

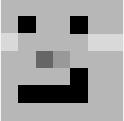
Write out image result

- **Language:** C++ implemented from scratch
- **Optimizations and programming model:**
  - OpenMP (Algorithm parallelism)
  - OpenMP (Data parallelism)
  - MPI+OpenMP+Slurm (Hybrid)
- **Platform:** Google Cloud (Virtual Instance, Cluster with Slurm)

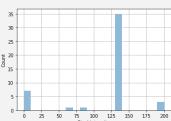
# Optimization 1: OpenMP - Algorithm parallelism

```
For cell_img in all images
{
```

Step 0: Input image



Step 1: Pixel intensity histogram



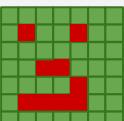
Parallelism  
using OpenMP

Step 2: Find threshold

```
0, 0, 0, 0, 0, 0, 0,  
60,  
.....  
87,  
130, 130, 130, 130, 130, 130...  
200, 200, 200
```

Parallelism  
using OpenMP

Step 3: Final segmentation

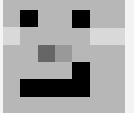


Parallelism  
using OpenMP

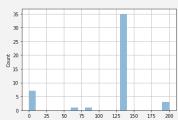
```
}
```

# Optimization 1: OpenMP - Algorithm parallelism

```
For cell_img in all images
```

```
{  
    Step 0: Input image  

```

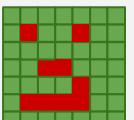
```
    Step 1: Pixel intensity histogram
```



```
    Step 2: Find threshold
```

```
    0, 0, 0, 0, 0, 0, 0,  
    60,  
    .....  
    87,  
    130, 130, 130, 130, 130, 130...  
    200, 200, 200
```

```
    Step 3: Final segmentation
```

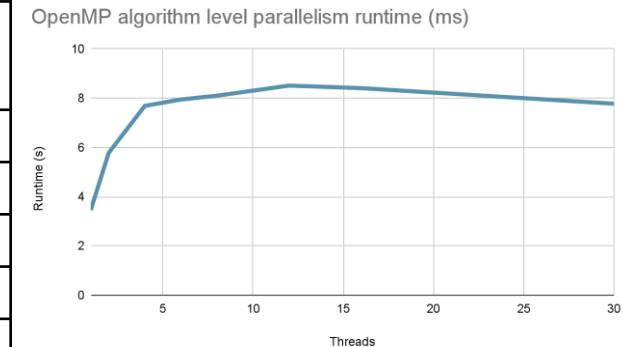


Parallelism using OpenMP

Parallelism using OpenMP

Parallelism using OpenMP

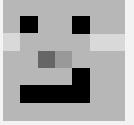
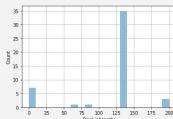
Threads	Runtime (s)	Speedup compared to 1 thread
1	3.456	1
2	5.779	0.59799
4	7.687	0.4496
6	7.945	0.4359
8	8.101	0.4266
12	8.510	0.40614
16	8.407	0.4111
30	7.775	0.4445

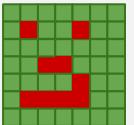


**Google Cloud:** Compute optimized VM with 30 vCPUs.

# Optimization 1: OpenMP - Algorithm parallelism

```
For cell_img in all images
```

- {
  - Step 0: Input image
  - Step 1: Pixel intensity histogram
  - Step 2: Find threshold

```
0, 0, 0, 0, 0, 0, 0, 0,  
60,  
.....  
87,  
130, 130, 130, 130, 130...  
200, 200, 200
```
  - Step 3: Final segmentation

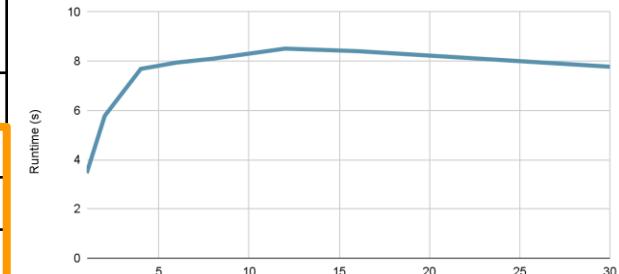
Parallelism using OpenMP

Parallelism using OpenMP

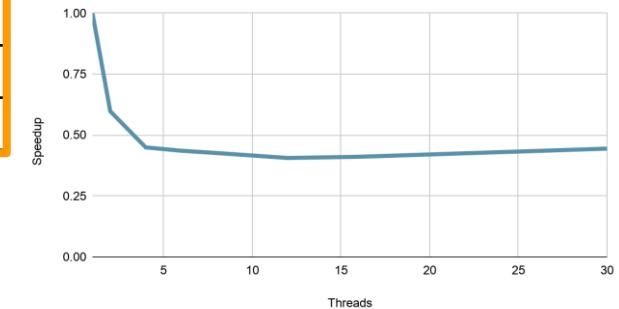
Parallelism using OpenMP

Threads	Runtime (s)	Speedup compared to 1 thread
1	3.456	1
2	5.779	0.59799
4	7.687	0.4496
6	7.945	0.4359
8	8.101	0.4266
12	8.510	0.40614
16	8.407	0.4111
30	7.775	0.4445

OpenMP algorithm level parallelism runtime (ms)



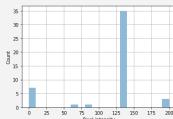
Speedup compared to 1 thread

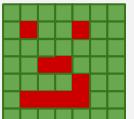


**Google Cloud:** Compute optimized VM with 30 vCPUs.

# Optimization 1: OpenMP - Algorithm parallelism

For cell\_img in all images

- {
  - Step 0: Input image
  - Step 1: Pixel intensity histogram
  - Step 2: Find threshold

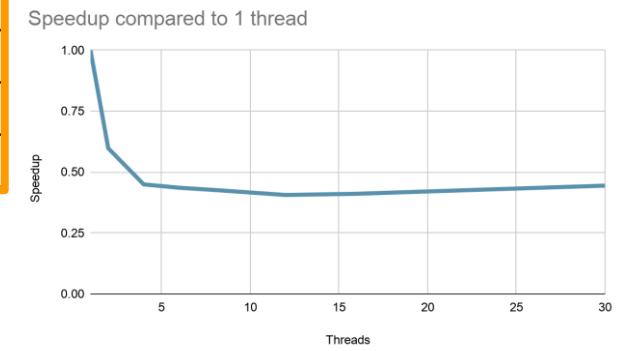
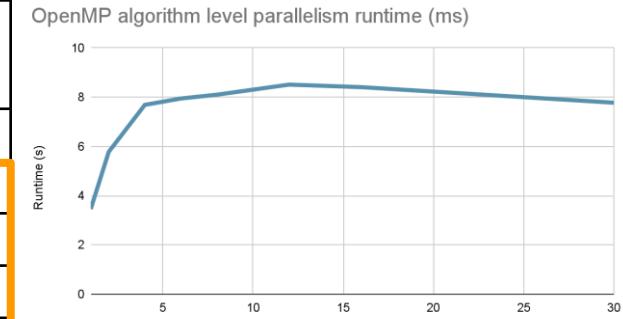
```
0, 0, 0, 0, 0, 0, 0, 0,  
60,  
.....  
87,  
130, 130, 130, 130, 130, 130...  
200, 200, 200
```
  - Step 3: Final segmentation

Parallelism using OpenMP

Parallelism using OpenMP

Parallelism using OpenMP

Threads	Runtime (s)	Speedup compared to 1 thread
1	3.456	1
2	5.779	0.59799
4	7.687	0.4496
6	7.945	0.4359
8	8.101	0.4266
12	8.510	0.40614
16	8.407	0.4111
30	7.775	0.4445



Amdahl's law theoretical max speedup 1.58.

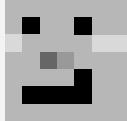
Since each of these parallelized subparts of not compute intensive, the overhead immediately dominates.

# Optimization 2: OpenMP - Data parallelism

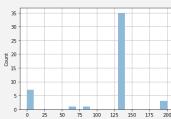
```
For cell_img in all images
```

```
{
```

Step 0: Input image



Step 1: Pixel intensity histogram



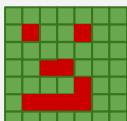
Parallelism  
using OpenMP

Step 2: Find threshold

```
0, 0, 0, 0, 0, 0,  
60,  
.....  
87,  
130, 130, 130, 130, 130...  
200, 200, 200
```

Parallelism  
using OpenMP

Step 3: Final segmentation



Parallelism  
using OpenMP

```
}
```

# pragma statements before the for loop

Allowing multiple cell images to be simultaneously segmented.

# Optimization 2: OpenMP - Data parallelism

Thread number	Total time(ms)	Speedup of current compared to previous	Speedup of current compared to 1 thread	Step 0: data loading (%)	Step 1: compute histogram (%)	Step 2: finding threshold (%)	Step 3: final segmentation (%)	Step 4: data save (%)
1	3.574s	NA	1	0.46368	0.208419	0.440318	0.19899	0.08334
2	1.794759	1.9915	1.9915	0.471731	0.219885	0.0152618	0.205544	0.0870849
4	0.976506s	1.8379	3.6602	0.473829	0.215001	0.0153819	0.206288	0.0890124
6	0.702265s	1.3905	5.0896	0.471621	0.216026	0.0152663	0.207052	0.089535
8	0.618021s	1.13631	5.7835	0.464095	0.222237	0.0157273	0.208233	0.089284
12	0.539055s	1.14648	6.6307	0.450517	0.218754	0.0158821	0.219383	0.0947447
16	0.442824s	1.2173	8.0716	0.454073	0.218009	0.0157286	0.215558	0.0959675
30	0.387223s	1.1436	<b>9.2306</b>	0.536978	0.180146	0.0124914	0.181916	0.0880761
60	0.393248s	0.9847	9.0892	0.760144	0.0795287	0.00555554	0.0804626	0.0733615
90	0.424771s	0.9258	8.4147	0.808752	0.0474346	0.0032952	0.0478846	0.0925408

# Optimization 3: MPI + OpenMP + Slurm on GCP

- **Google Cloud Cluster:**
  - 8 nodes
  - Each with 4 vCPUs
  - With SLURM
- **Hybrid Model:**
  - MPI
  - OpenMP
- **Parallel Types:**
  - Data parallelism
  - Algorithm parallelism
- **Best speedup:**
  - **5.096**

node\thread	1 threads	2 threads	4 threads	8 threads
1 nodes	<b>6.570s</b>	5.490s	4.785s	4.594s
2 nodes	3.292s	2.709s	2.525s	2.557s
4 nodes	2.211s	1.726s	1.880s	1.848s
6 nodes	1.628s	<b>1.435s</b>	1.403s	1.389s
8 nodes	1.432s	<b>1.369s</b>	1.320s	<b>1.289s</b>

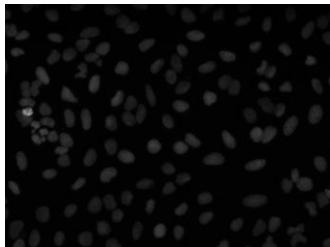
# Results



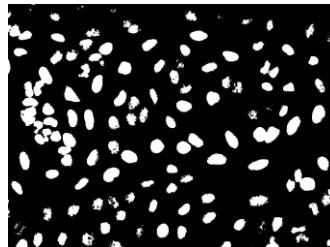
# Results and Discussion

## Example 1:

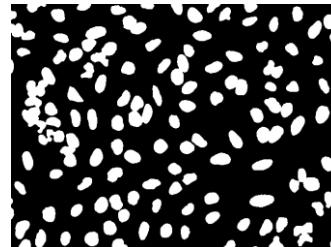
Original image



Otsu's result

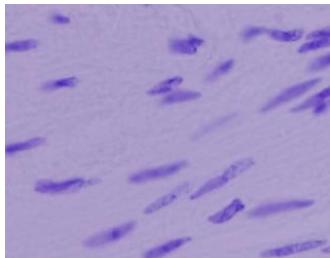


Ground truth



## Example 2:

Original image



Otsu's result



Ground truth



### Speed:

- Extremely fast in speed
- Great speed up with a maximum of **9.23**

### Accuracy:

- Performs well in most cell segmentations with pixel intensity contrast

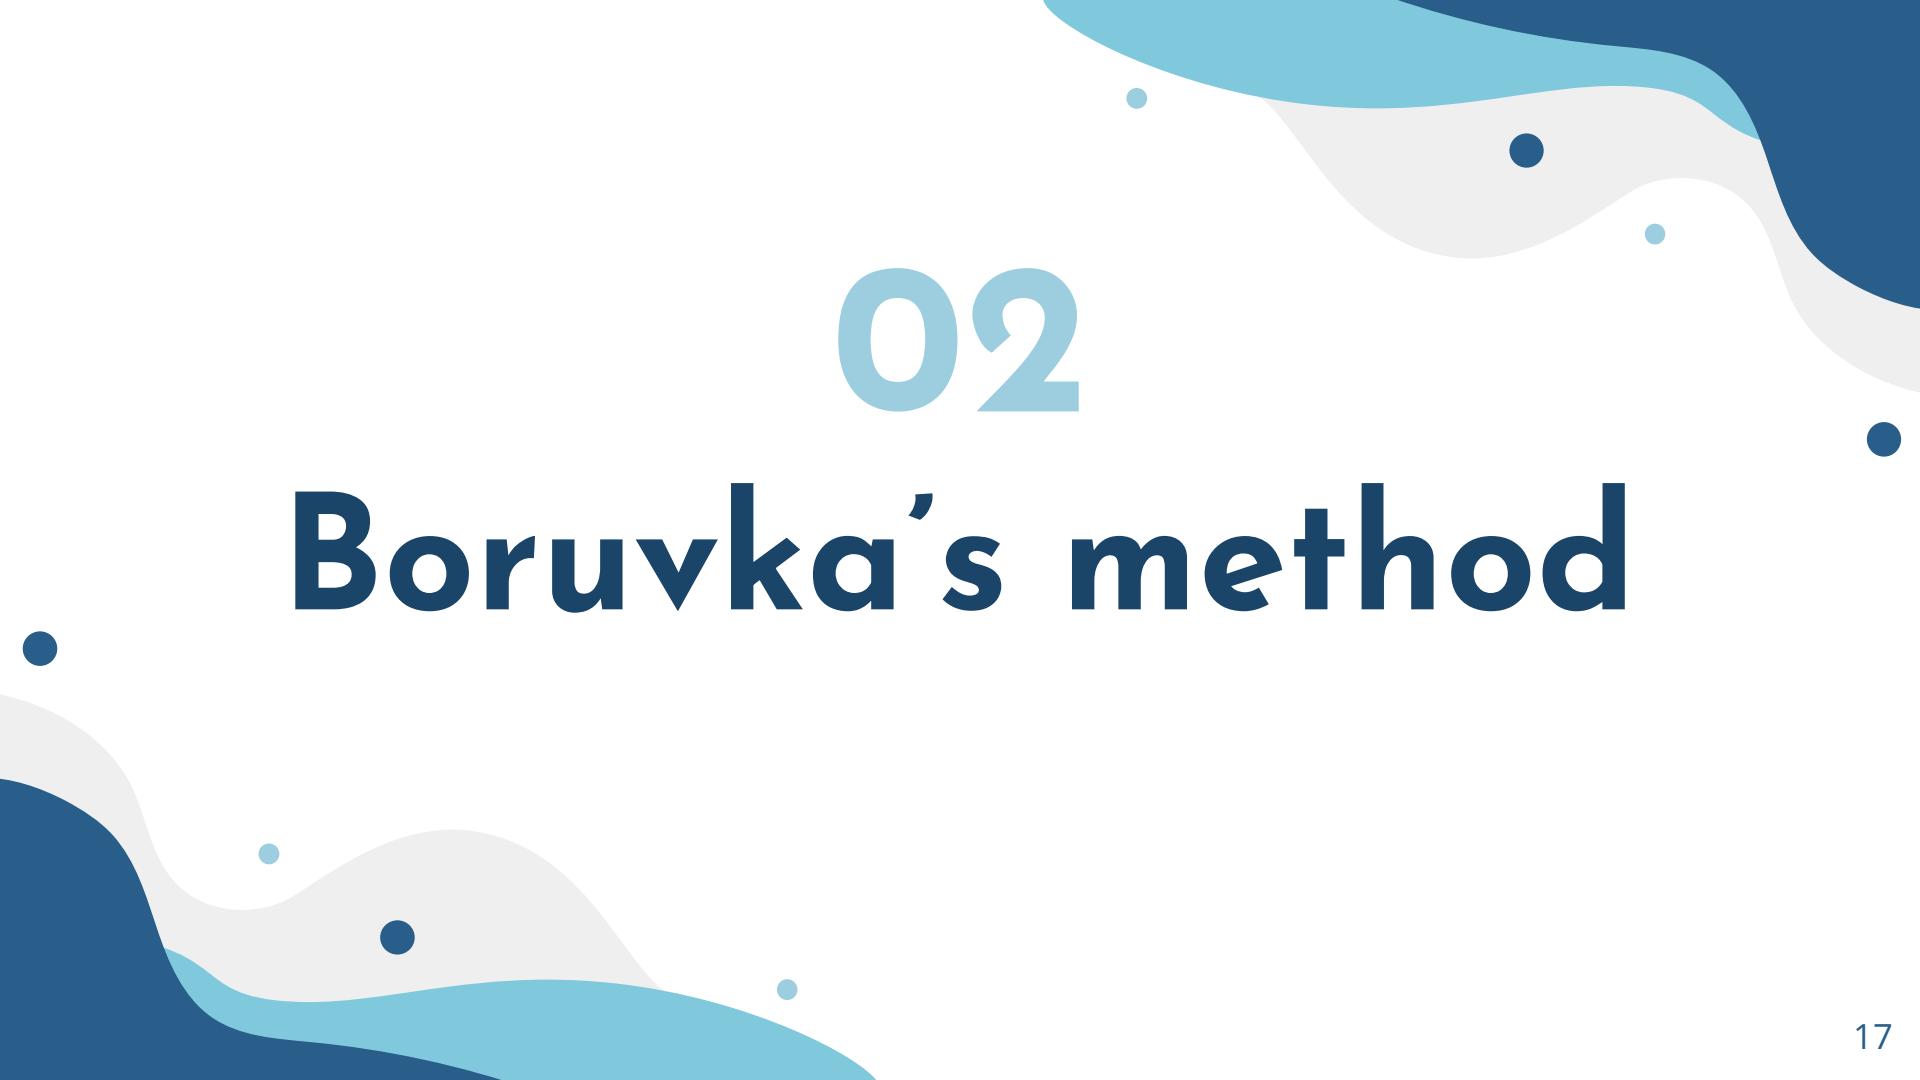
### Pros:

- Fast and mostly accurate
- Little dependency
- Doesn't need training

### Cons:

- Unable to differentiate foreground and background

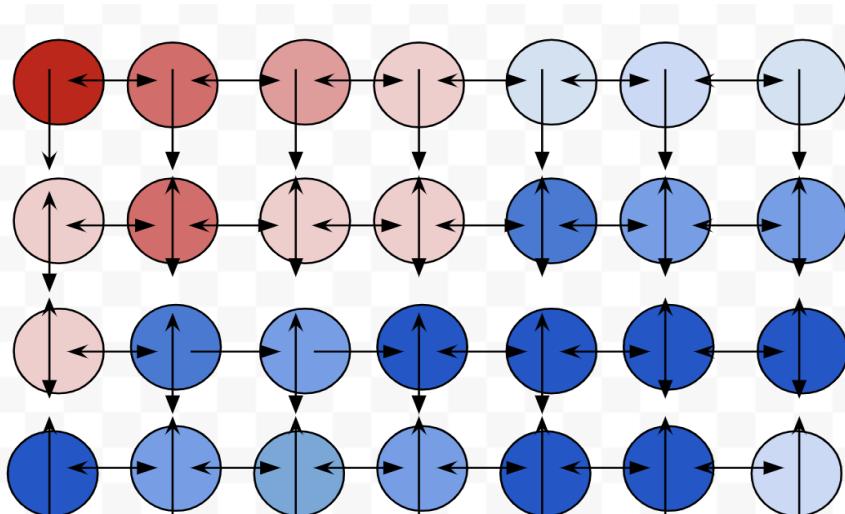
**Not compute intensive:** fast, but not friendly for parallelism



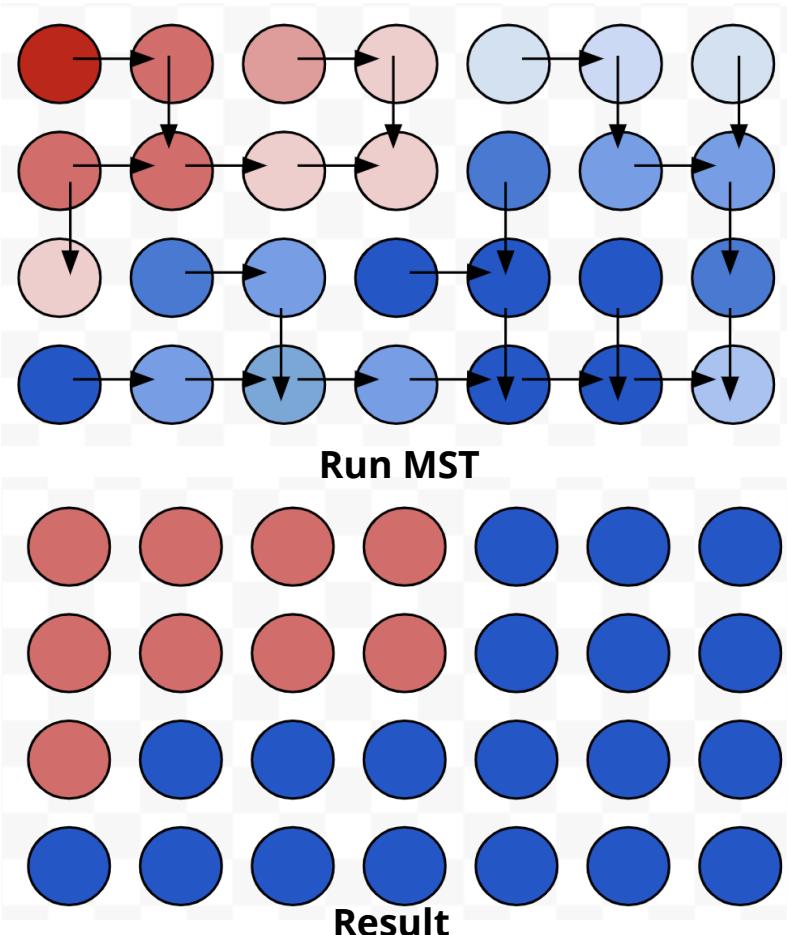
02

# Boruvka's method

# Algorithm



Initialize

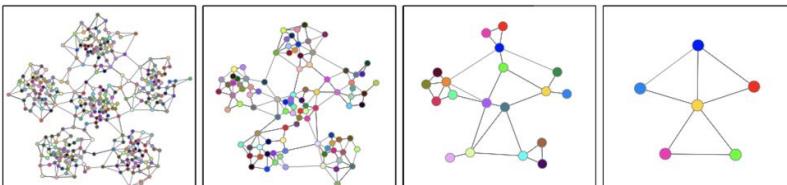


# Optimization 1: Memory

- Overheads: Storing Graph

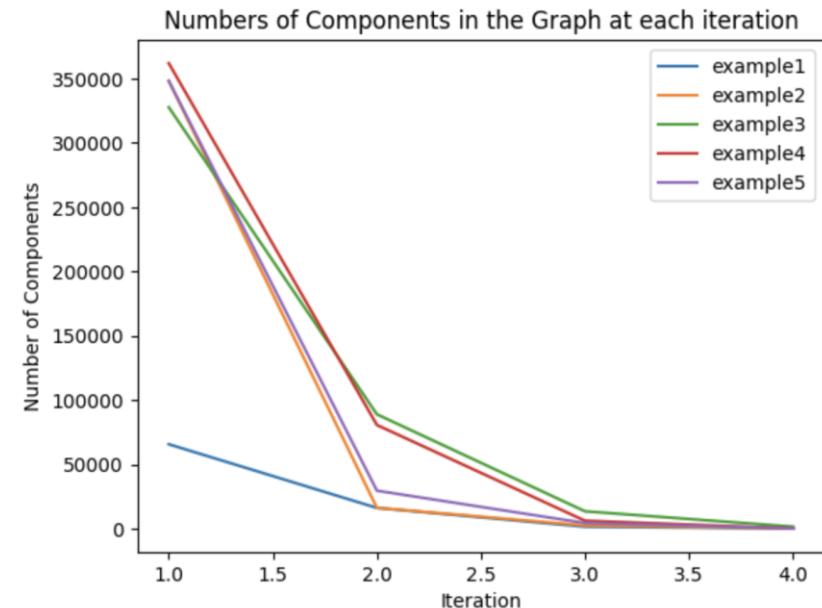
3.5 hours

per 256\*256 images



## Solution

- Edge compaction



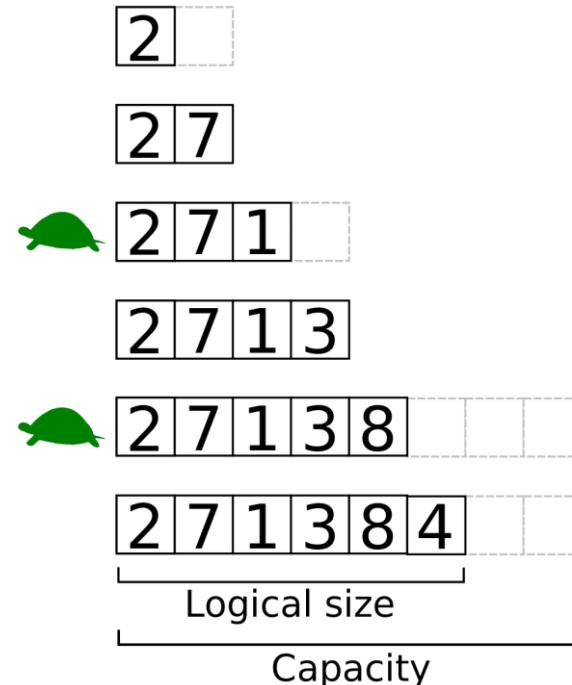
# Optimization 1: Memory

## Solution (cont.)

- Passed by reference, NOT pass by value!

(Note: In C++, things are passed by value unless you specify otherwise)

- C++ dynamic memory allocation & deallocation, e.g. malloc(), calloc(), free() delete[] ...



C++'s `std::vector`: dynamic array

# Optimization 1: Memory

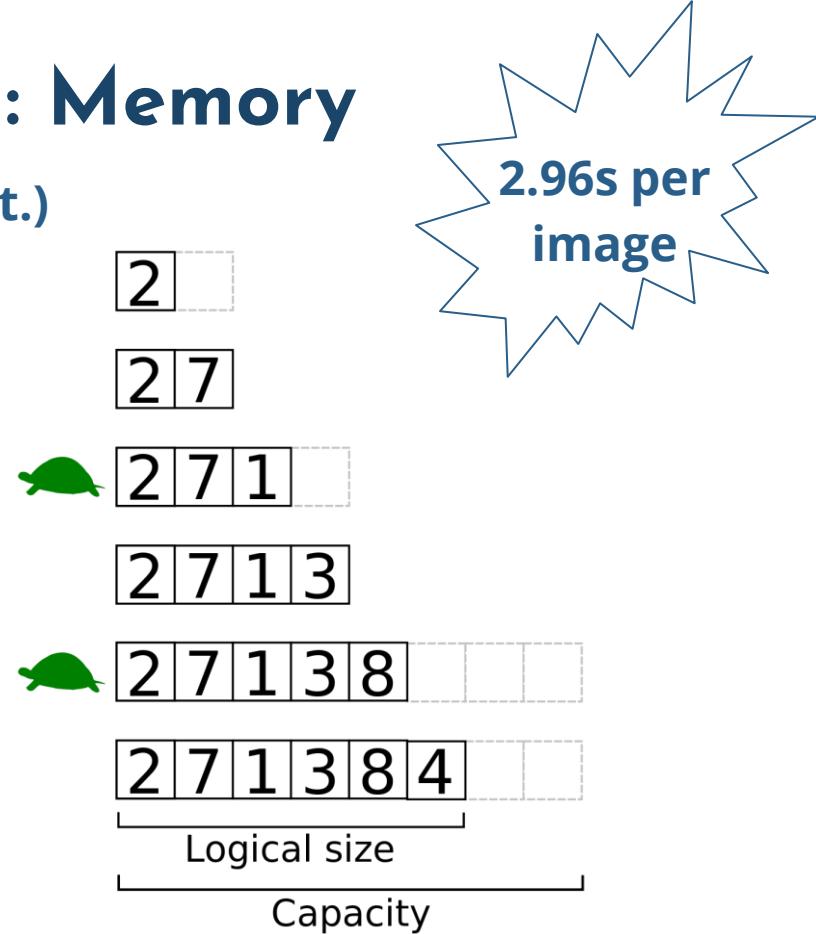
## Solution (cont.)

**Don't write code like this!**

```
std::vector<int> nums;  
for(int i=0; i<n; i++){  
    nums.push_back(1);  
}
```

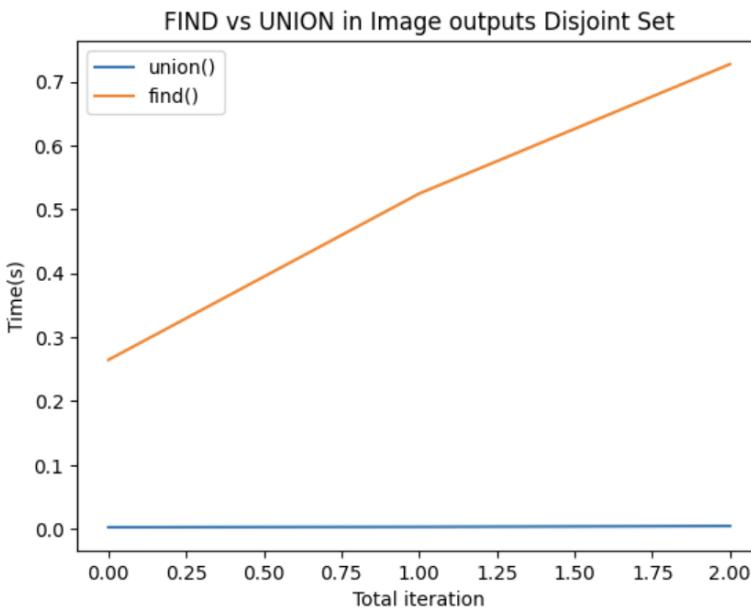
**Instead,**

```
std::vector<int> nums(n, 1);
```



# Optimization 2: OpenMP

- Now, it takes 2.96s per image!  
But processing all the images  
still takes 33m7.4s!



Total  
9m 30s

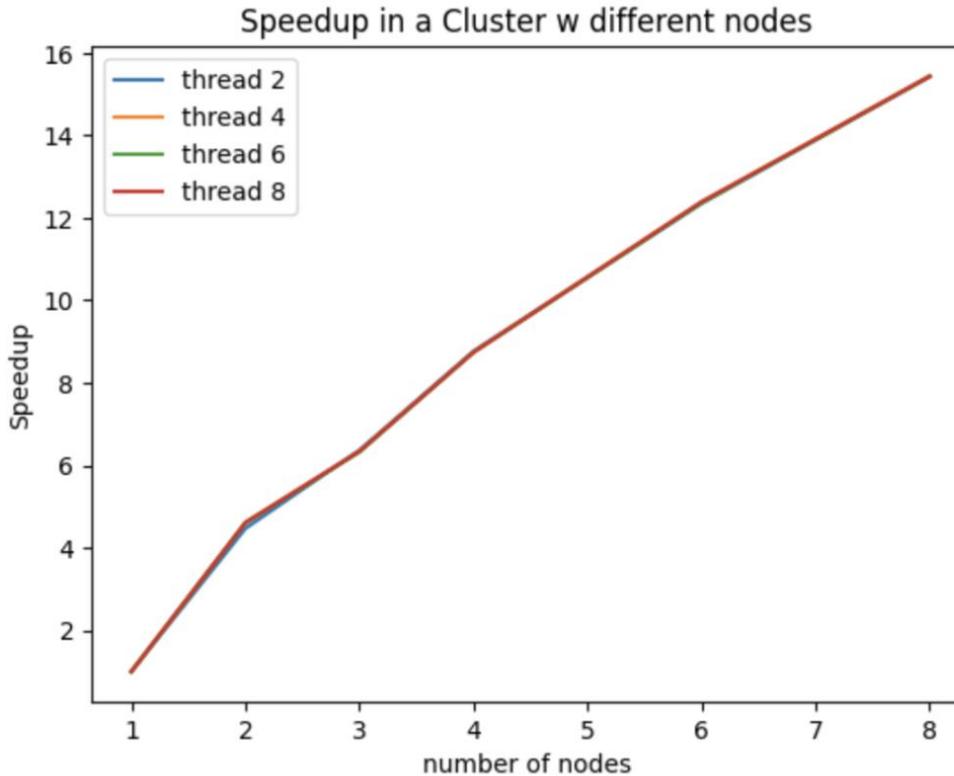
## Parallelize find() operation

- OpenMP, #pragma omp parallel for schedule(dynamic)

## Overhead: Data Dependency

- C++ built-in function for atomic memory access:  
`_sync_bool_compare_and_swap`

# Optimization 3: MPI & OpenMP



1 node 8 thread/node, 30vCPU

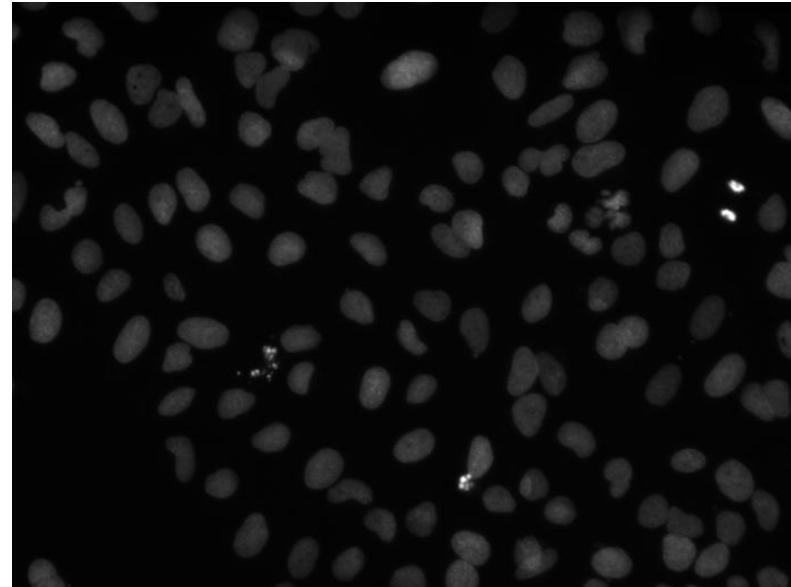
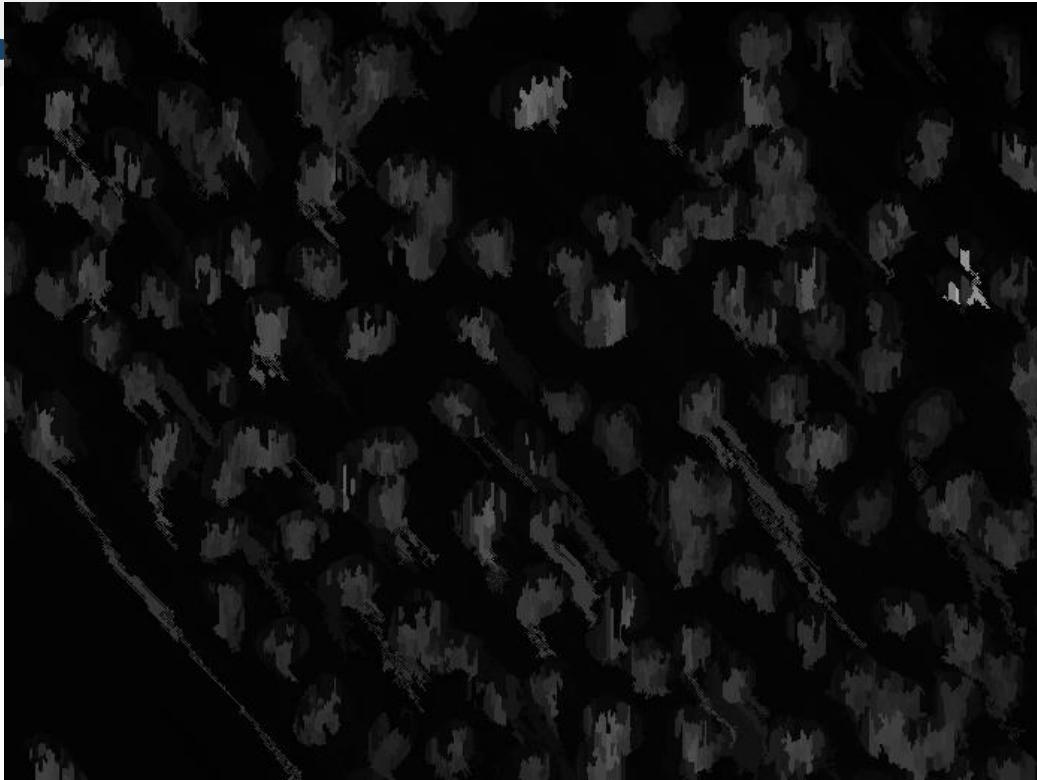
**9m 30s**

⟨Data Parallelism⟩

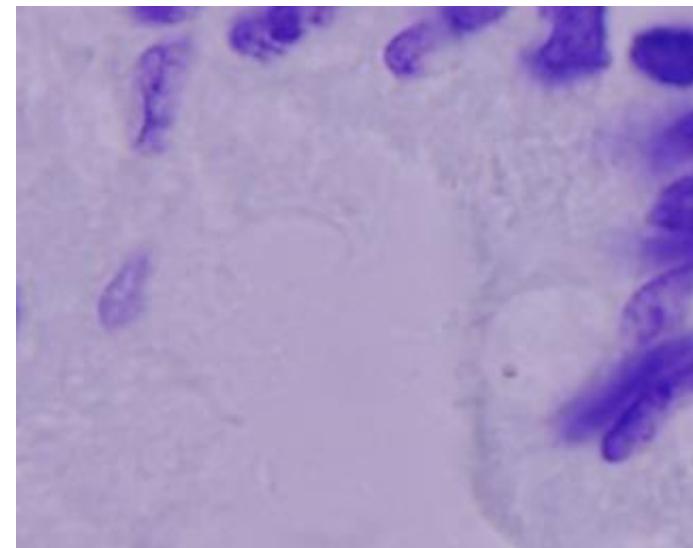
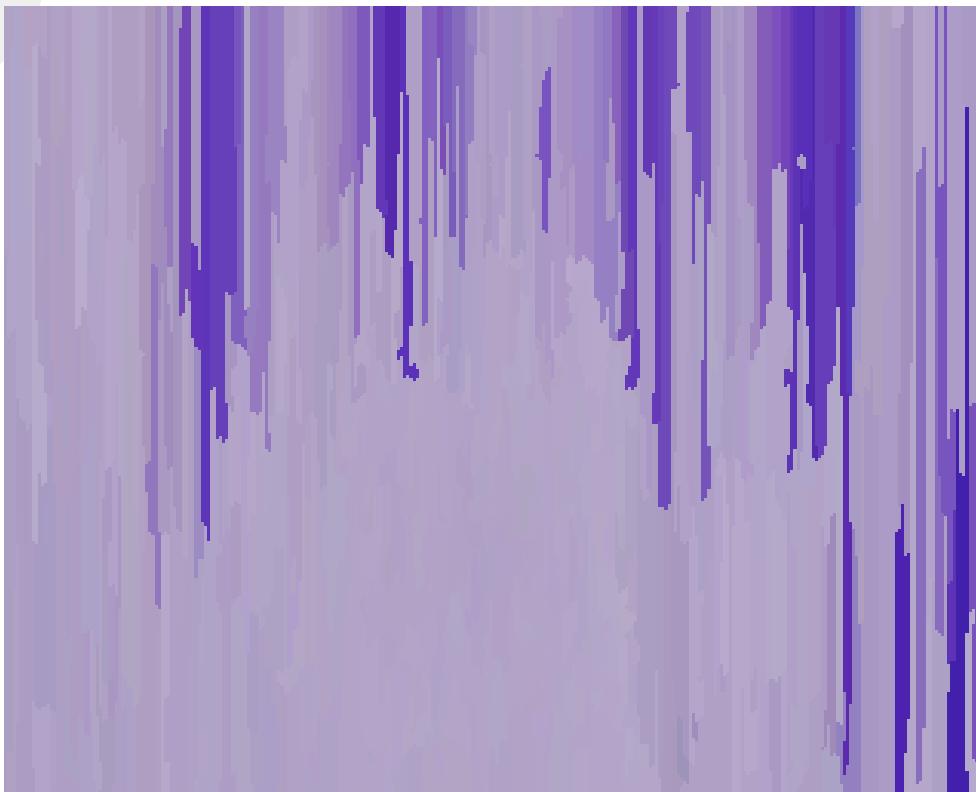
8 nodes 4 threads/node, 4vCPU

**2m 8.79s**

# Good Results



# Bad



**Boruvka: When there is a big color difference, it can do great!**

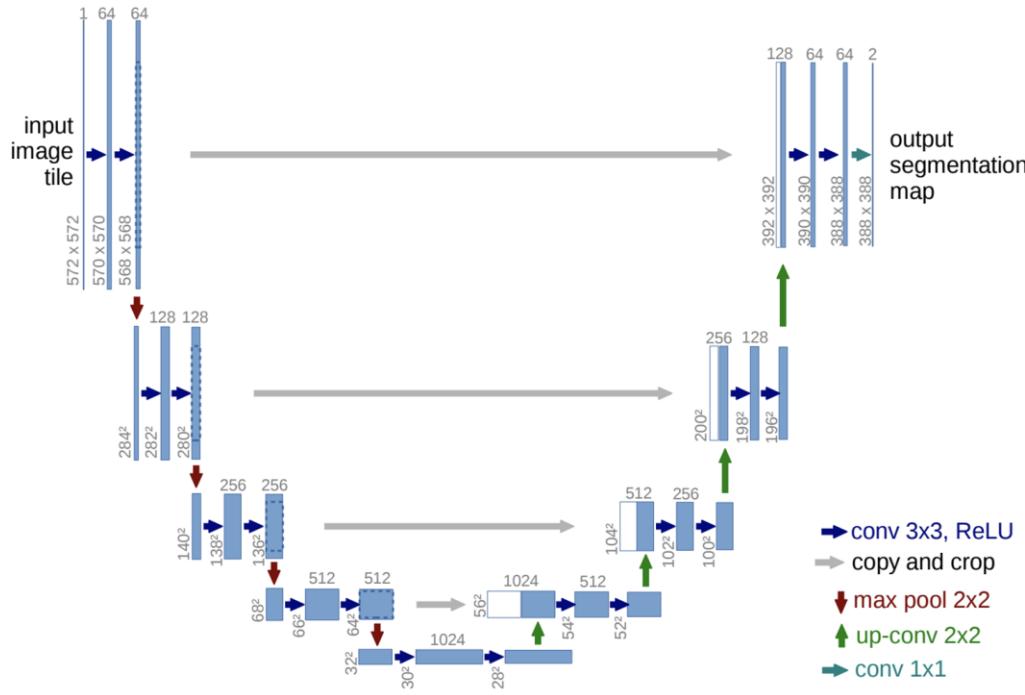




# 03

# UNet

# UNet: Model Overview



## Convolutional Neural Network

- Highly non-linear function mapping each pixel to 0 or 1
- Huge amount of model parameters
  - Total params: 1,941,105
  - Trainable params: 1,941,105
  - Non-trainable params: 0
- Learns fast with data augmentation

# UNet

## Programming Model and Platform

### GPU accelerated computing

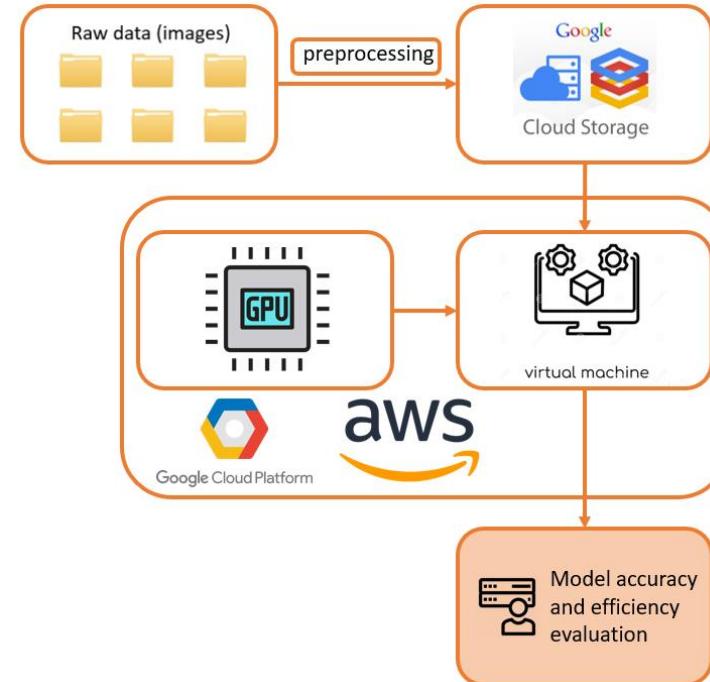
- Single GPU of different types
- Parallel data augmentation and training

### Software:

- Python3
- Keras + Tensorflow

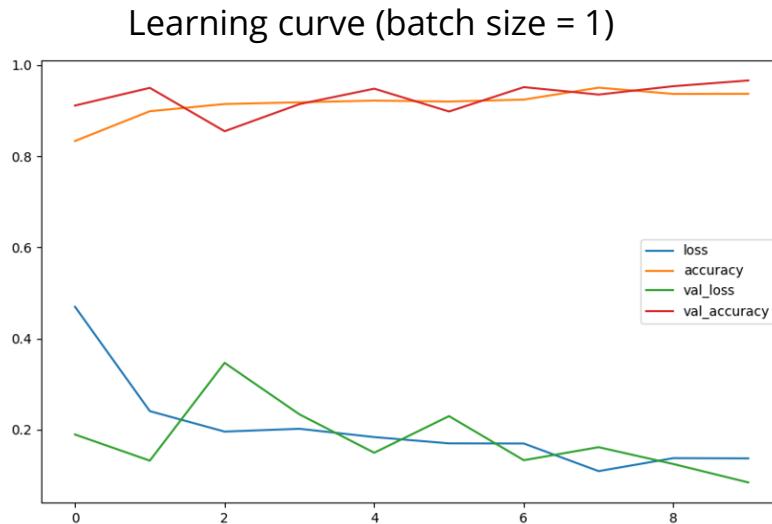
### Platform:

- Google Cloud Platform (GCP)
- AWS

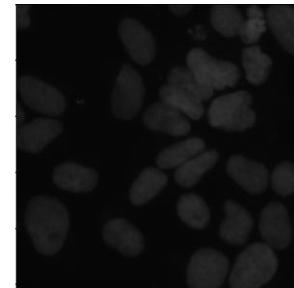


# Segmentation Results

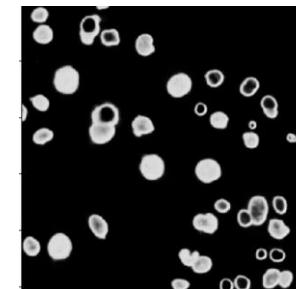
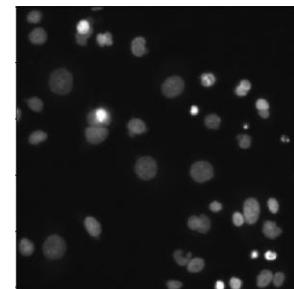
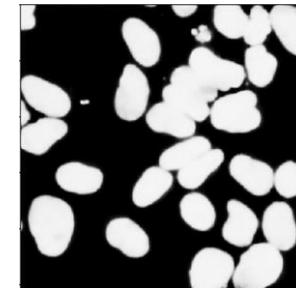
Final pixel classification accuracy: 95%



Input Image



UNet Output



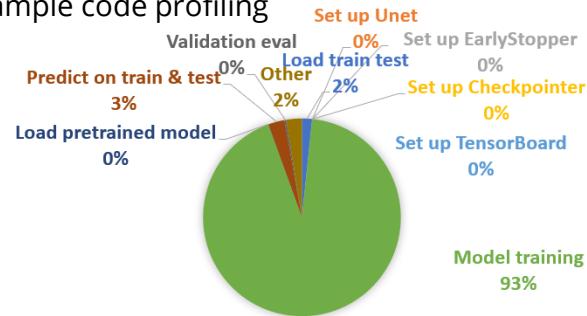
# GPU Acceleration

Six VMs (including the baseline)

Instance		GCP c2-standard-30	AWS g3.4xlarge	AWS p2.xlarge	AWS p3.2xlarge	AWS g4dn.4xlarge	GCP n1-standard-16
GPU info	Type	N/A	Tesla M60	Tesla K80	Tesla V100 SXM2	Tesla T4	Tesla P4
	Memory	N/A	7618MiB	11441MiB	16160MiB	15109MiB	7611MiB
	Count	0	1	1	1	1	1
CPU info	Type	Intel(R) Xeon(R) CPU	Intel(R) Xeon(R) Platinum 8259CL @2.50GHz	Intel(R) Xeon(R) E5-2686 v4 @2.30GHz	Intel(R) Xeon(R) E5-2686 v4 @2.30GHz	Intel(R) Xeon(R) Platinum 8259CL @2.50GHz	Intel(R) Xeon(R) CPU @2.30GHz
	Count	30	16	4	8	16	16
	Thread(s)/core	2	2	2	2	2	2
	Core(s)/socket	15	8	2	4	8	8
	Socket(s)	1	1	1	1	1	1
	Memory	120 GiB	122 GiB	61 GiB	61 GiB	64 GiB	64 GiB
	Price (instance only)	\$1.5660/hr	\$1.14/hr	\$0.90/hr	\$3.06/hr	\$1.204/hr	\$1.376944/hr
Batchsize = 1 Training speed		138s/epoch 696ms/step	14s/epoch 68ms/step	30s/epoch 152ms/step	14s/epoch 70ms/step	12s/epoch 60ms/step	13s/epoch 67ms/step

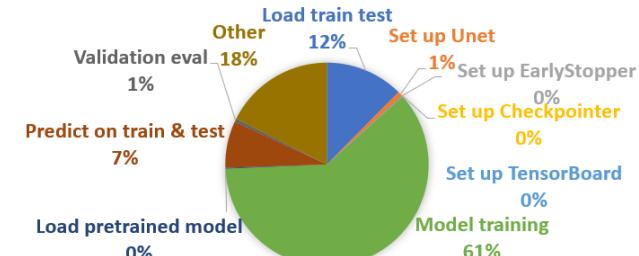
\* Much more detailed report can be found in the write-up.

Sample code profiling



BASELINE RUNTIME BREAKDOWN

Script total time:  
**1521.629 sec**

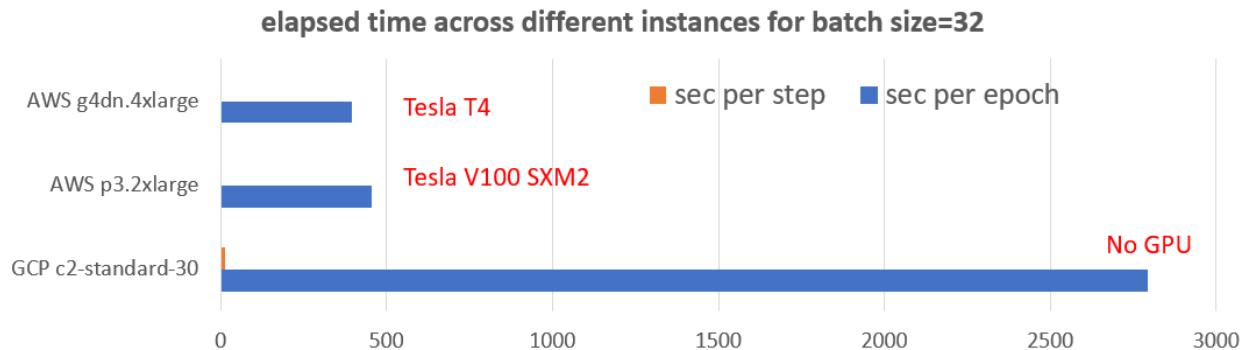
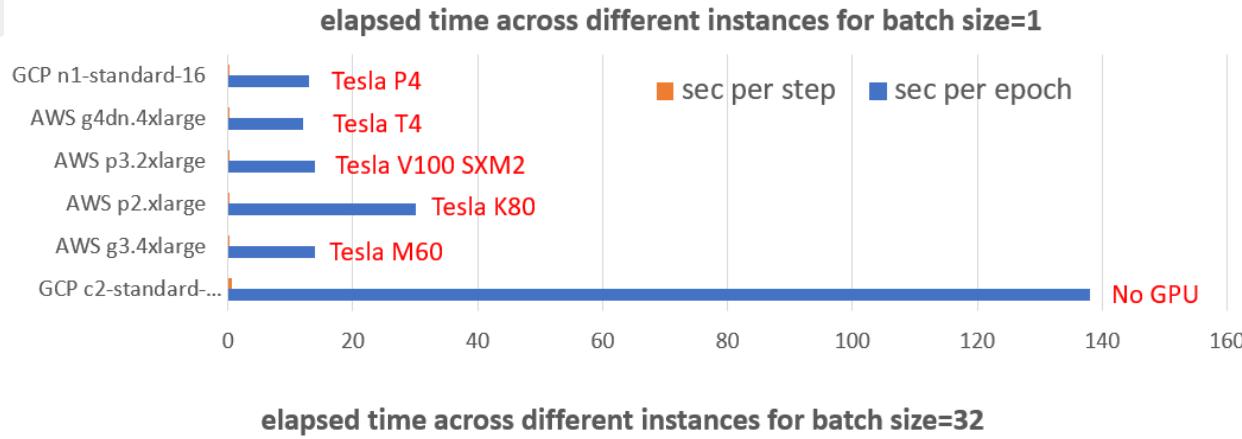


GCP GPU RUNTIME  
BREAKDOWN

Script total time:  
**245.146 sec**

\* Also used cProfile, TensorBoard. Please refer to the write-up.

# Speedup and Scalability

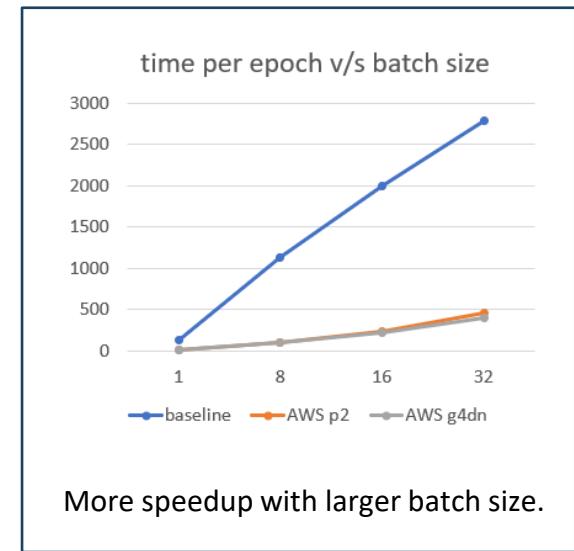


**Best speedup:** AWS g4dn.4xlarge

~ 11 for batch size = 1

~ 7 for batch size = 32

\* compared to baseline

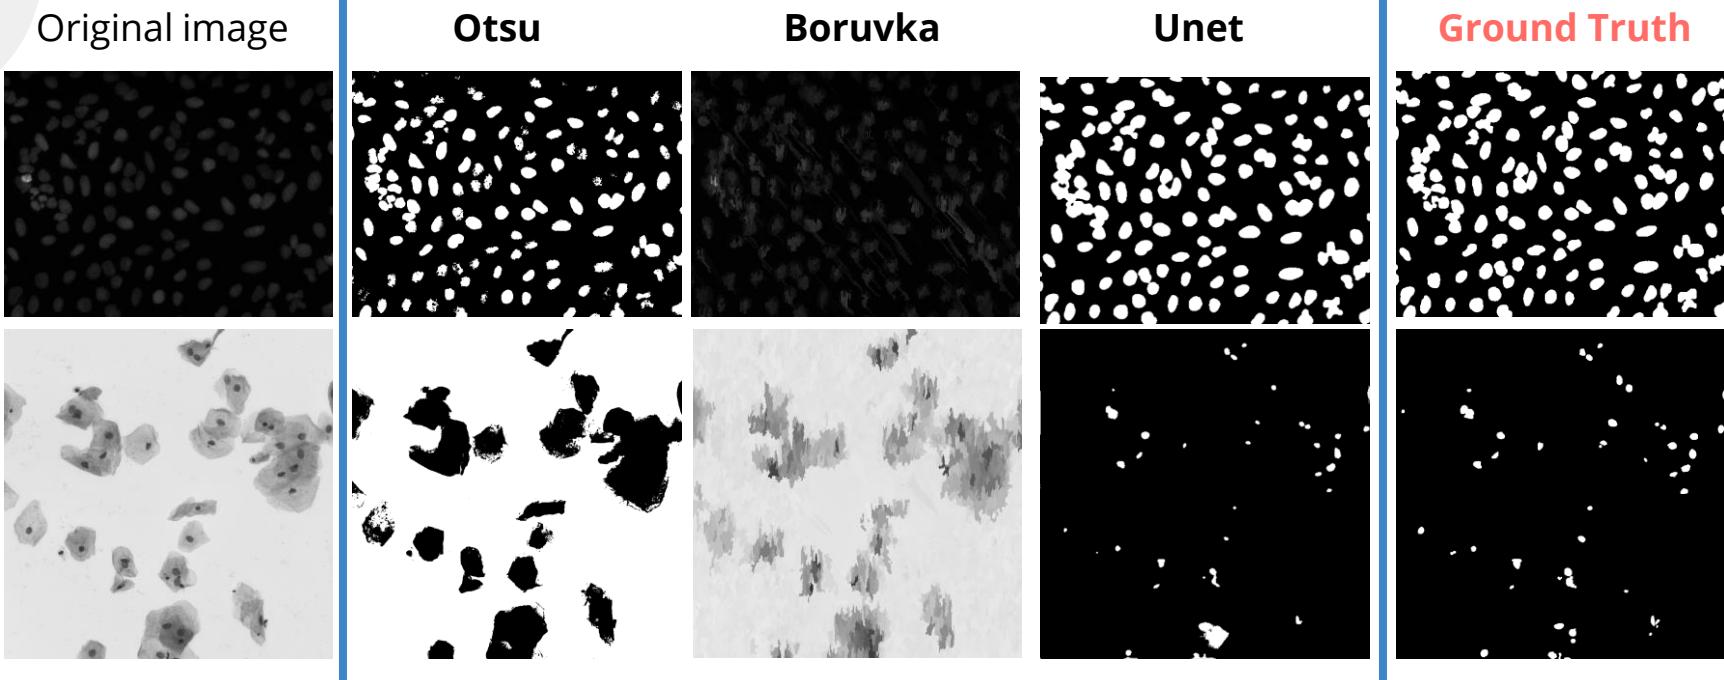




04

# Final Recap

# Key Results Recap



# • Levels of Parallelism/Cloud Infra Used

## Otsu's / Boruvka's Algorithm

Optimization:

- MPI
- OpenMP

Cloud infrastructure:

- Google Cloud VM
- Google Cloud Cluster

Service:

- Slurm

Levels of Parallelism:

- Task level
- Procedure level
- Loop level

Types of Parallelism:

- Inter node: Data
- Intra node: Function

Language: C++

## UNet

Optimization:

- GPU
- Spark

Cloud infrastructure:

- Google Cloud VM + GPU
- AWS VM + GPU

Service:

- Google Cloud Dataproc Cluster

Levels of Parallelism:

- Task level

Types of Parallelism:

- Data parallelism

Language: Python3

# Future Work

In addition to the significant speedup of our optimized implementations, we can further leverage

- Multi-GPU instances
- Split huge images for parallel processing

# Thank you for watching!

Special thanks to **Haipeng** for discussion and suggestions